

# benchmarking and profiling rkumaraswamy1/2

Sebastian Fischer

11/04/2020

In this file the two rng-functions for the kumaraswamy distribution are profiled:

- rkumaraswamy1: inverse transform sampling, utilizing `runif(0,1)`
- rkumaraswamy2: sampling from `rbeta(1, b)` and calculating  $x^{1/a}$

The mark function from the bench package was used to benchmark the functions. The profvis function from the provis package was used to profile the functions.

The input checking and the ‘actual calculations’, i.e. the `rkumaraswamy1_main` and `rkumaraswamy2_main`, are benchmarked separately. As both rng-functions use the same input-checking, there is no point in comparing them.

## Benchmarking

One thing to notice is that the mean and the max of the runtime are not included intentionally. The reason is that the benchmarking process can be interrupted by other programs. Therefore some runtimes are higher than they would be without these interruptions. This makes the mean and the max less informative.

### Benchmarking the input checking

The input checking is benchmarked separately for scalar parameters and for zed iparameters. For both  $n$  was set to  $1e5$

For the scalar input checking the following values were chosen:

$$a = 0.5, \quad b = 0.3, \quad \min = -1, \quad \max = 2$$

For the vectorized input checking 100 values from a standard normal were generated for each of the parameters  $a, b, \min, \max$  (naturally before the benchmarking).

The results are summarized in the following table:

input	min	median	itr/sec	mem_alloc	gc/sec	n_itr	n_gc	total_time
scalar	$205\mu s$	$242\mu s$	3405	$172KB$	8.47	1609	4	$473ms$
vectorized	$6.04ms$	$8.69ms$	89.52	$8.78MB$	39.39	25	11	$279ms$

The important point to take away from this table is that the time, required to check the inputs, as well as the memory usage is negligible. (One side note is that the input checking for a scalar parameter does not scale with  $n$  but for vectorized parameters it does because of the recycling.)

## Benchmarking the ‘actual calculations’

The following table compares the two functions: `rkumaraswamy1_main` and `rkumaraswamy2_main`. As the results for vectorized parameters and scalar parameters are similar, only the result for scalar parameters is provided.  $n$  was again set to 100 000.

function	min	median	itr/sec	mem_alloc	gc/sec	n_itr	n_gc	total_time
<code>rkumaraswamy1</code>	20.9ms	23ms	41.23	3.54MB	14.73	14	5	340ms
<code>rkumaraswamy2</code>	34ms	35.4ms	26.51	5.5MB	4.82	11	2	415ms

The first implementation outperforms the second version. This is not really surprising however, as Wikipedia states that

“It (Kumaraswamy distribution) is similar to the Beta distribution, but much simpler to use especially in simulation studies since its probability density function, cumulative distribution function and quantile functions can be expressed in closed form.”

## Profiling

Only `rkumaraswamy1_main` and `rkumaraswamy2_main` are profiled, as they are in the main interest.  $n$  is set to  $2e6$ .

Flame Graph	Data	Options ▾		
Code	File	Memory (MB)		Time (ms)
▼ <code>rkumaraswamy2_main</code>	<expr>	-106.8	91.6	760
<code>rbeta</code>	<code>kumaraswamy-rng.R</code>	0	30.5	700
<code>x[!invalid_par] &lt;- (rbeta(n = n, shape1 = 1, shape2 = b))[!invalid_par]</code>	<code>kumaraswamy-rng.R</code>	-61.0	22.9	30
<code>x[!invalid_par] &lt;- (x ^ (1 / a) * (max - min) + min)[!invalid_par]</code>	<code>kumaraswamy-rng.R</code>	-45.8	38.1	30
▼ <code>rkumaraswamy1_main</code>	<expr>	0	38.1	470
<code>qkumaraswamy_main</code>	<code>kumaraswamy-rng.R</code>	0	30.5	350
<code>runif</code>	<code>kumaraswamy-rng.R</code>	0	0	110
<code>p[!invalid_par] &lt;- runif(n_valid)</code>	<code>kumaraswamy-rng.R</code>	0	7.6	10
Sample Interval: 10ms		1230ms		

Figure 1: profiling data for `rkumaraswamy1/2_main`

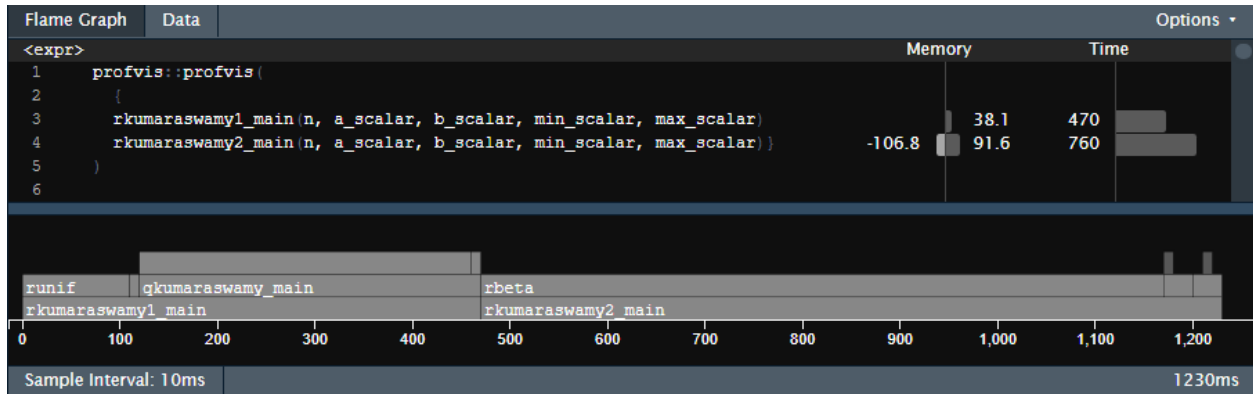


Figure 2: profiling time-line for `rkumaraswamy1/2_main`

The results confirm that indeed the generation of pseudo-random numbers from the beta distribution is the bottleneck in `rkumaraswamy2_main` and the reason why it is so much slower than the first version. In the

paper, on the basis of which the `rbeta` function was programmed (see the help-page of `rbeta`), it says that it is based on a rejection method. This explains, why it is so much slower compared to `runif`. Not only does it have to sample twice: once from the proposal and once from the uniform distribution, the number of times it has to sample from these distributions is potentially much higher than  $n$  if the rejection rate is high. This naturally also explains the far higher memory usage, as more samples are generated during the runtime.

For the `rkumaraswamy2_main` function, the most costly operation is the quantile function. Even though it is analytically calculable, it still requires 9 vectorized operations to calculate the quantiles for the values generated by `runif`.